

# Deep Learning

Natural Language Processing (NLP)

Lionel Fillatre

April 2021

# Outline

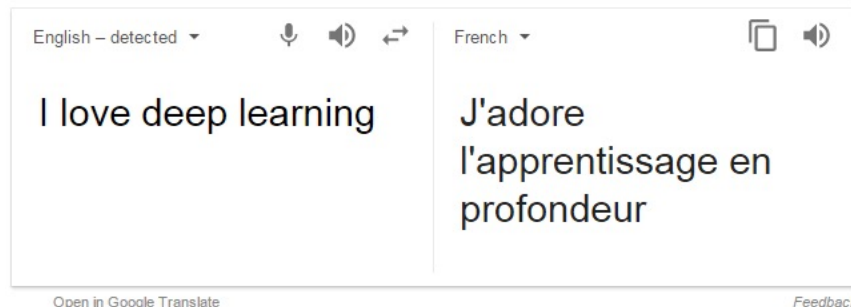
- Introduction
- Word Embedding
- Word2Vec
- Conclusion

# Introduction

# Natural Language Processing (NLP) Applications

- Sentiment Analysis
- Email Filters
- Voice Recognition
- Information Extraction
- Translation
- ...

	Sentence	Class	index
0	So there is no way for me to plug it in here i...	0	0
1	Good case, Excellent value.	1	1
2	Great for the jawbone.	1	2
3	Tied to charger for conversations lasting more...	0	3
4	The mic is great.	1	4
5	I have to jiggle the plug to get it to line up...	0	5
6	If you have several dozen or several hundred c...	0	6
7	If you are Razr owner...you must have this!	1	7
8	Needless to say, I wasted my money.	0	8
9	What a waste of money and time!	0	9
10	And the sound quality is great.	1	10



# Language Model

# Language Models


- Formal grammars (e.g. regular, context free) give a hard “binary” model of the legal sentences in a language.
- For NLP, a probabilistic model of a language that gives a probability that a string is a member of a language is more useful.
- To specify a correct probability distribution, the probability of all sentences in a language must sum to 1.

# Uses of Language Models

- Speech recognition
  - “I ate a cherry” is a more likely sentence than “Eye eight uh Jerry”
- OCR & Handwriting recognition
  - More probable sentences are more likely correct readings.
- Machine translation
  - More likely sentences are probably better translations.
- Generation
  - More likely sentences are probably better Natural Language generations.
- Context sensitive spelling correction
  - “Their are problems wit this sentence.”

# Language Modeling

- Language Modeling is the task of predicting what word comes next



- Given a sequence of words  $w_1, w_2, \dots, w_t$ , compute the probability distribution of the next word  $w_{t+1}$

$$\Pr(w_{t+1} | w_t, w_{t-1}, \dots, w_1)$$

where  $w_{t+1}$  can be any word in the vocabulary  $V = \{v_1, v_2, \dots, v_{|V|}\}$

- A system that does this is called a Language Model.



# Completion Prediction

- A language model also supports predicting the completion of a sentence.
  - Please turn off your cell \_\_\_\_\_
  - Your program does not \_\_\_\_\_
- Predictive text input systems can guess what you are typing and give choices on how to complete it.

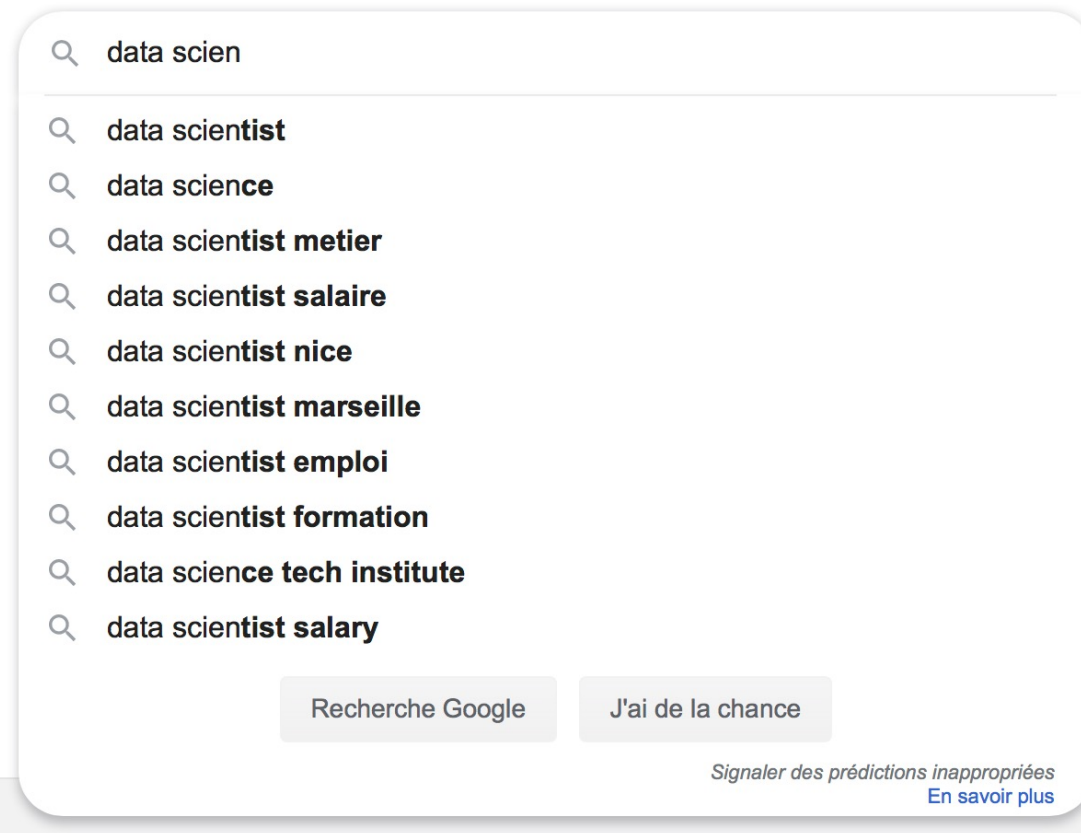
# Language Modeling

- You can also think of a Language Model as a system that assigns probability to a piece of text.
- For example, if we have some text  $w_1, w_2, \dots, w_T$ , then the probability of this text (according to the Language Model) is:

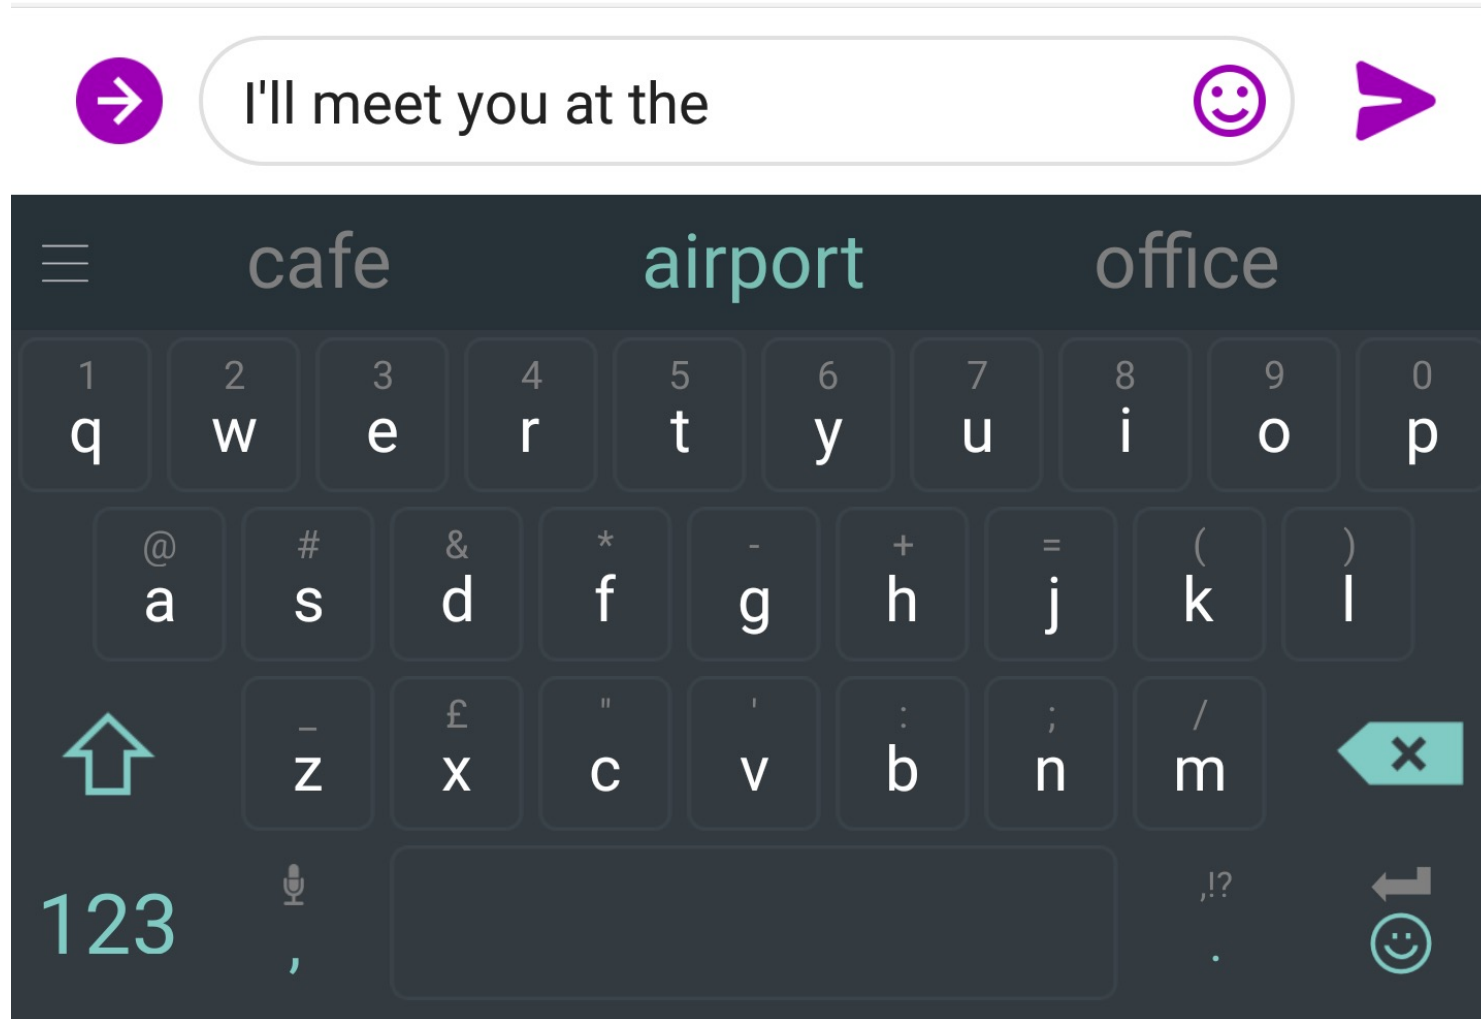
$$\Pr(w_1, w_2, \dots, w_T) = \prod_{t=1}^T \Pr(w_t | w_{t-1}, \dots, w_1)$$

- Example:  $\Pr(\text{its water was so transparent}) = \Pr(\text{its}) * \Pr(\text{water} | \text{its}) * \Pr(\text{was} | \text{its water}) * \Pr(\text{so} | \text{its water was}) * \Pr(\text{transparent} | \text{its water was so})$
- Language Modeling provides  $\Pr(w_t | w_{t-1}, \dots, w_1)$

# You use Language Models every day!



# You use Language Models every day!



# N-Gram

# N-gram Language Models

- Question: How to learn a Language Model?
- Answer (pre-Deep Learning): Learn a  $n$ -gram Language Model!
- Definition: A  $n$ -gram is a chunk of  $n$  consecutive words.
  - unigrams: “the”, “students”, “opened”, “their”
  - bigrams: “the students”, “students opened”, “opened their”
  - trigrams: “the students opened”, “students opened their”
  - 4-grams: “the students opened their”
- Idea: Collect statistics about how frequent different  $n$ -grams are, and use these to predict next word.

# N-Gram Models

- Estimate probability of each word given prior context.
  - $P(\text{phone} \mid \text{Please turn off your cell})$
- Number of parameters required grows exponentially with the number of words of prior context.
- An  $N$ -gram model uses only  $N-1$  words of prior context.
  - Unigram:  $P(\text{phone})$
  - Bigram:  $P(\text{phone} \mid \text{cell})$
  - Trigram:  $P(\text{phone} \mid \text{your cell})$
- Markov model
  - The Markov assumption is the presumption that the future behavior of a dynamical system only depends on its recent history.
  - In particular, in a  $k$ th-order Markov model, the next state only depends on the  $k$  most recent states, therefore an  $N$ -gram model is a  $(N-1)$ -order Markov model.

# N-gram Language Models

- First we make a simplifying assumption: preceding  $n - 1$  words

$$\Pr(w_{t+1} | w_t, \dots, w_1) = \Pr(w_{t+1} | w_t, \dots, w_{t-n+2})$$

- Conditional probabilities:

$$\Pr(w_1 | w_t, \dots, w_{t-n+2}) = \frac{\Pr(w_{t+1}, w_t, \dots, w_{t-n+2})}{\Pr(w_t, \dots, w_{t-n+2})}$$

- **Question:** How do we get these  $n$ -gram and  $(n - 1)$ -gram probabilities?
- **Answer:** By counting them in some large corpus of text!
  - Statistical approximation:

$$\Pr(w_{t+1} | w_t, \dots, w_{t-n+2}) \approx \frac{\text{count}(w_{t+1}, w_t, \dots, w_{t-n+2})}{\text{count}(w_t, \dots, w_{t-n+2})}$$



# Estimating Probabilities

- $N$ -gram conditional probabilities can be estimated from raw text based on the relative frequency of word sequences
  - Bigram:  $\Pr(w_t | w_{t-1}) = \frac{\text{count}(w_{t-1}w_t)}{\text{count}(w_{t-1})}$
  - N-gram:  $\Pr(w_t | w_{t-N+1}^{t-1}) = \frac{\text{count}(w_{t-N+1}^{t-1}w_t)}{\text{count}(w_{t-N+1}^{t-1})}$
- To have a consistent probabilistic model, append a unique start (<s>) and end (</s>) symbol to every sentence and treat these as additional words, e.g.,
  - <s> I am Sam </s>
  - <s> Sam I am </s>

# An Example

- $\langle s \rangle$  I am Sam  $\langle /s \rangle$
- $\langle s \rangle$  Sam I am  $\langle /s \rangle$
- $\langle s \rangle$  I do not like green eggs and ham  $\langle /s \rangle$

$$\begin{array}{lll} P(I | \langle s \rangle) = \frac{2}{3} = .67 & P(\text{Sam} | \langle s \rangle) = \frac{1}{3} = .33 & P(\text{am} | I) = \frac{2}{3} = .67 \\ P(\langle /s \rangle | \text{Sam}) = \frac{1}{2} = 0.5 & P(\text{Sam} | \text{am}) = \frac{1}{2} = .5 & P(\text{do} | I) = \frac{1}{3} = .33 \end{array}$$

# $N$ -gram Language Models: Example

- Suppose we are learning a 4-gram Language Model

~~as the proctor started the clock, the students opened their~~ \_\_\_\_\_

$$\Pr(w | \text{students opened their}) = \frac{\text{count}(\text{students opened their } w)}{\text{count}(\text{students opened their})}$$

- For example, suppose that in the corpus:
  - “students opened their” occurred 1000 times
  - “students opened their books” occurred 400 times  
→  $\Pr(\text{books} | \text{students opened their}) = 0.4$
  - “students opened their exams” occurred 100 times  
→  $\Pr(\text{exams} | \text{students opened their}) = 0.1$

# Sparsity Problems with $n$ -gram Language Models

Problem: What if “students opened their  $w$ ” never occurred in data? Then  $w$  has probability 0!

(Partial) Solution: Add small  $\delta$  to the count for every  $w \in V$ . This is called smoothing.

$$\Pr(w|\text{students opened their}) = \frac{\text{count}(\text{students opened their } w)}{\text{count}(\text{students opened their})}$$


Problem: What if “students opened their” never occurred in data? Then we can not calculate probability for any  $w$  !

(Partial) Solution: Just condition on “opened their” instead. This is called backoff.

Note: Increasing  $n$  makes sparsity problems worse. Typically we can't have  $n$  bigger than 5.

# Storage Problems with $n$ -gram Language Models

Storage: Need to store count for all  $n$ -grams you saw in the corpus.


$$\Pr(w|\text{students opened their}) = \frac{\text{count}(\text{students opened their } w)}{\text{count}(\text{students opened their})}$$

Increasing  $n$  or increasing corpus increases model size!

# $n$ -gram Language Models in practice

- You can build a simple trigram Language Model over a 1.7 million word corpus (Reuters) in a few seconds on your laptop
- Example:
  - *today the* \_\_\_\_\_
  - The probability distribution is

Company	0.153
Bank	0.153
price	0.077
italian	0.039
emirate	0.039

- It seems reasonable but sparsity problem: not much granularity in the probability distribution to decide between « Company » and « Bank »

# Generating text with a n-gram Language Model

- You can use a Language Model to generate text.

*today the price of gold per ton , while production of shoe lasts and shoe industry , the bank intervened just after it considered and rejected an imf demand to rebuild depleted european stocks , sept 30 end primary 76 cts a share.*

- Surprisingly grammatical!
- ...but incoherent. We need to consider more than three words at a time if we want to model language well.
- But increasing  $n$  worsens sparsity problem, and increases model size...

# Word Embedding



# How do we represent the meaning of a word?

- Definition: meaning
  - The idea that is represented by a word, phrase, etc.
  - The idea that a person wants to express by using words, signs, etc.
  - The idea that is expressed in a work of writing, art, etc.
- Commonest linguistic way of thinking of meaning:

signifier (symbol)  $\Leftrightarrow$  signified (idea or thing)

# Representing words as discrete symbols

- In traditional NLP, we regard words as discrete symbols:
  - hotel, conference, motel – a localist representation
- Words can be represented by one-hot vectors:
  - motel = [0 0 0 0 0 0 0 0 0 0 1 0 0 0 0]
  - hotel = [0 0 0 0 0 0 0 1 0 0 0 0 0 0 0]
- One-hot encoding
  - Vector dimension = number of words in vocabulary  $V$  (e.g., 500,000)
  - The number of words in the vocabulary is  $|V|$
  - The vector consists of 0s in all cells with the exception of a single 1 in a cell used uniquely to identify the word.

# Problem with words as discrete symbols

- **Example:** in web search, if user searches for “Seattle motel”, we would like to match documents containing “Seattle hotel”.
- But:
  - motel = [0 0 0 0 0 0 0 0 0 0 1 0 0 0 0]
  - hotel = [0 0 0 0 0 0 0 1 0 0 0 0 0 0 0]
  - These two vectors are orthogonal.
- There is no natural notion of **similarity** for one-hot vectors!
- **Solution:**
  - Learn to encode similarity in the vectors themselves

# Distributional Hypothesis

- The Distributional Hypothesis in linguistics is derived from the semantic theory of language usage, i.e.
  - Words that are used and occur in the same contexts tend to purport similar meanings.
  - « A word is characterized by the company it keeps ». (J. R. Firth, 1957)
  - « The complete meaning of a word is always contextual, and no study of meaning apart from context can be taken seriously. »
- Use the many contexts of  $w$  to build up a representation of  $w$
- Example:
  - These context words will represent banking

*...government debt problems turning into **banking** crises as happened in 2009...*

*...saying that Europe needs unified **banking** regulation to replace the hodgepodge...*

*...India has just given its **banking** system a shot in the arm...*

# Word vectors


- We will build a dense vector for each word, chosen so that it is similar to vectors of words that appear in similar contexts

$$\textit{banking} = \begin{pmatrix} 0.286 \\ 0.792 \\ -0.177 \\ -0.107 \\ 0.109 \\ -0.542 \\ 0.349 \\ 0.271 \end{pmatrix}$$

- Note: word vectors are sometimes called word embeddings or word representations. They are a distributed representation.

# Representation Space

- Example 2D word embedding space, where similar words are found in similar locations



# Word2Vec

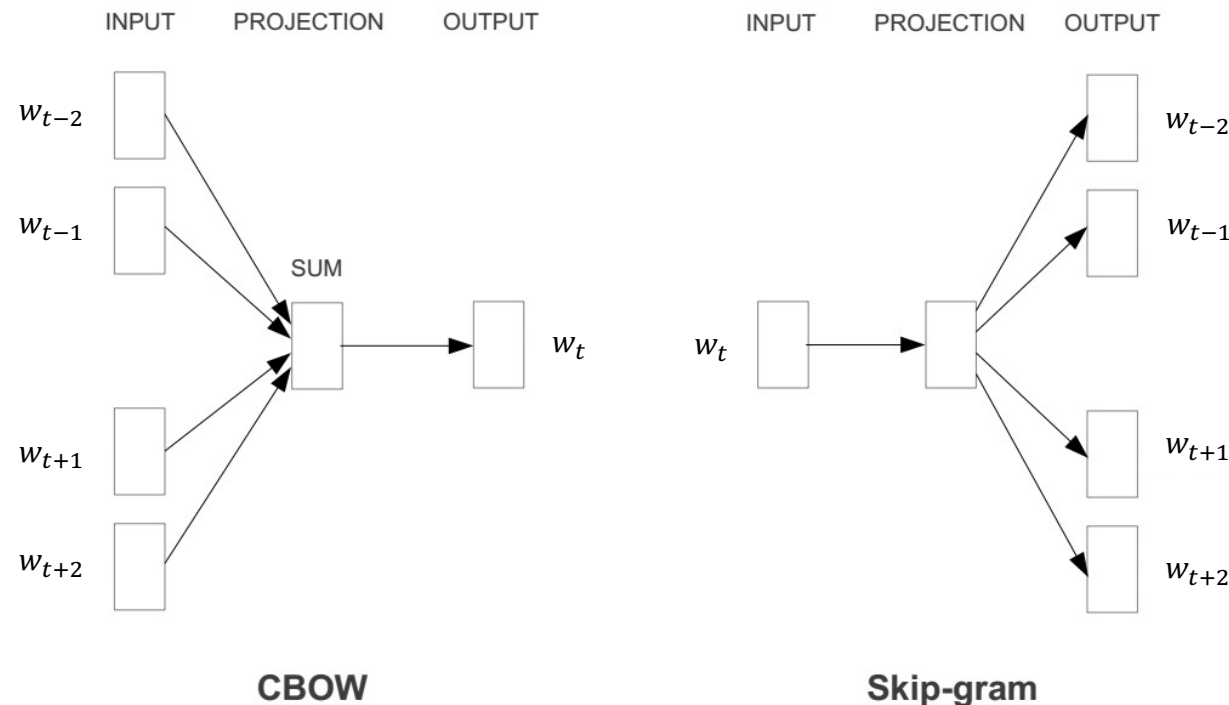
# Word2vec: Overview

- Word2vec (Mikolov et al. 2013) is a framework for learning word vectors
- Idea:
  - We have a large corpus of text
  - Every word in a fixed vocabulary is represented by a vector
  - Go through each position  $t$  in the text, which has a center word  $c$  and context (“outside”) words  $o$
  - Use the similarity of the word vectors for  $c$  and  $o$  to calculate the probability of  $o$  given  $c$  (or vice versa)
  - Keep adjusting the word vectors to maximize this probability



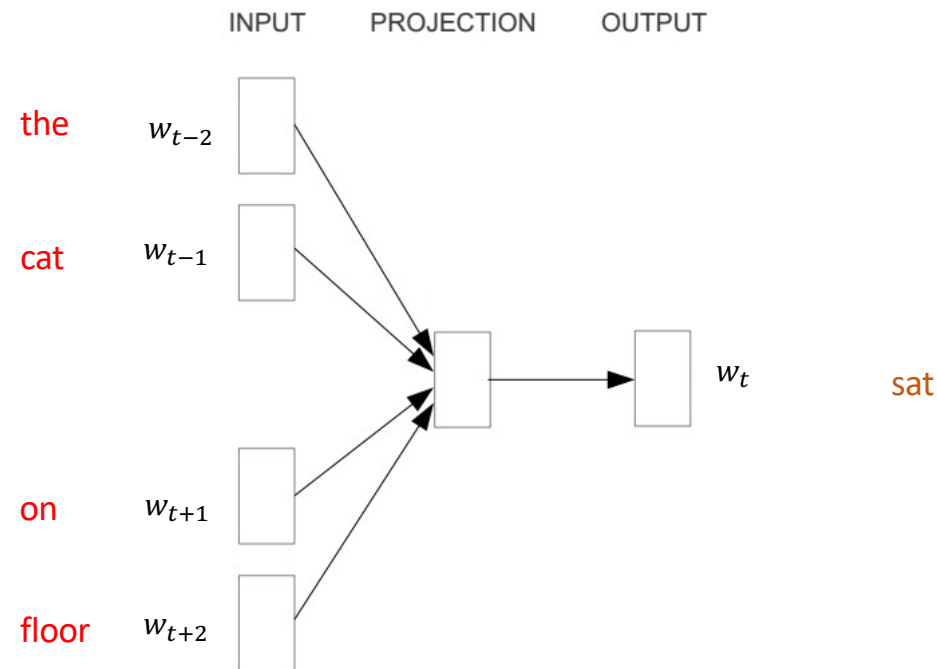
# Represent the meaning of word – word2vec

- 2 basic neural network models:
  - Continuous Bag of Word (CBOW): use a window of word to predict the middle word
  - Skip-gram (SG): use a word to predict the surrounding ones in window.




# Word2vec – Continuous Bag of Word


- E.g. “The cat sat on floor”
  - Window size = 2




# Input Encoding (window size = 1)



# Matrix Encoding




# Example




Note that we can take  $\hat{v} = \frac{v_{cat} + v_{on}}{2}$  instead of  $\hat{v} = v_{cat} + v_{on}$

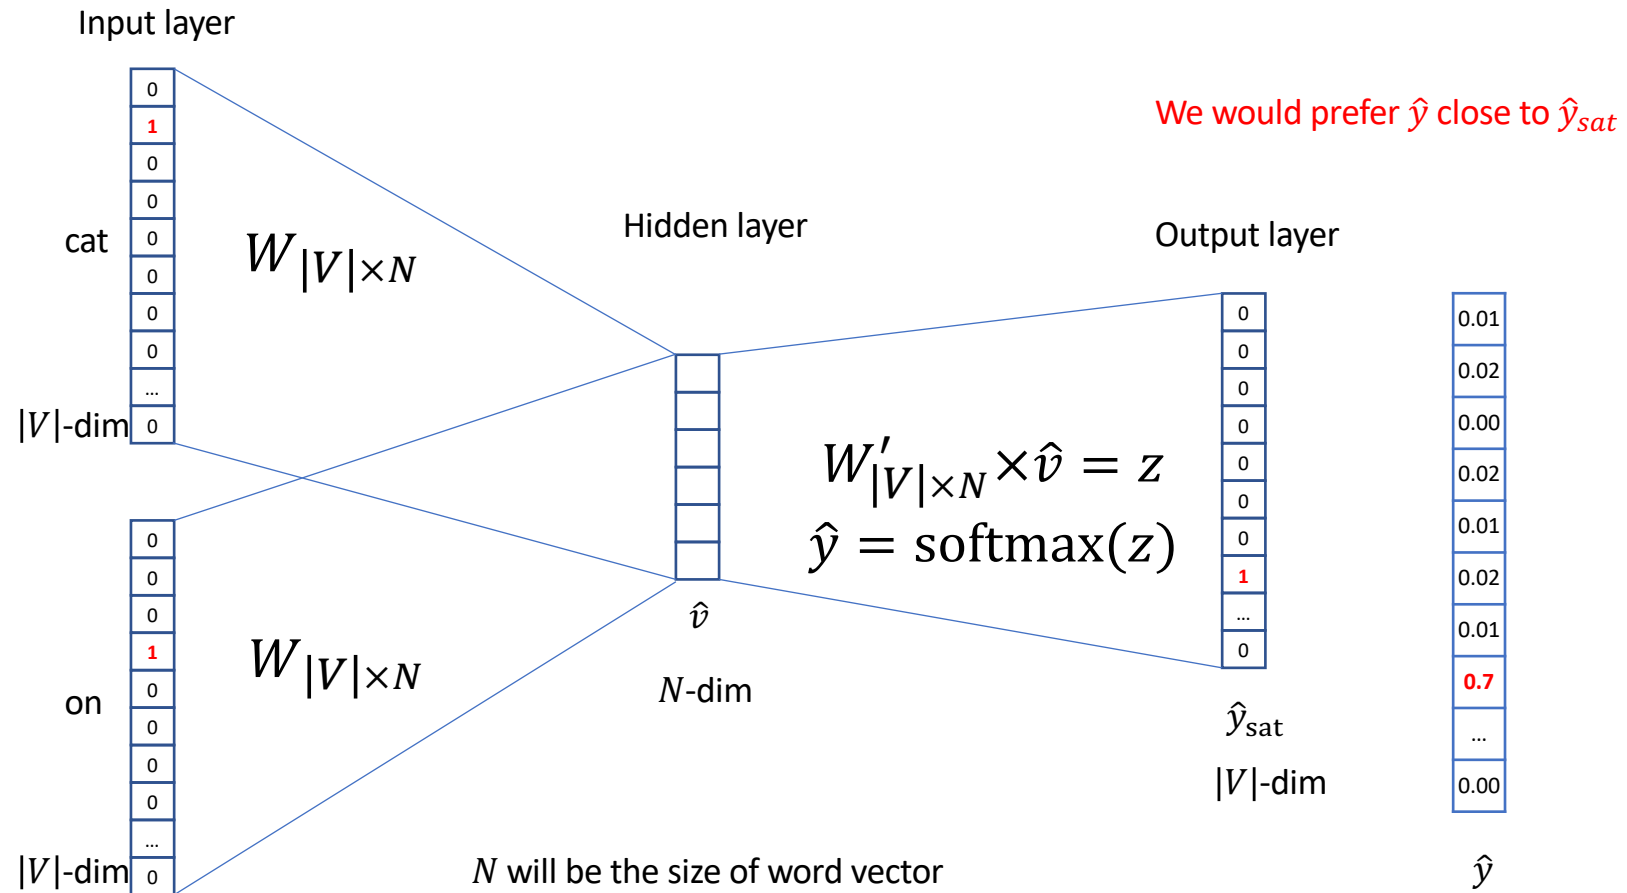
# Example



# Output Layer

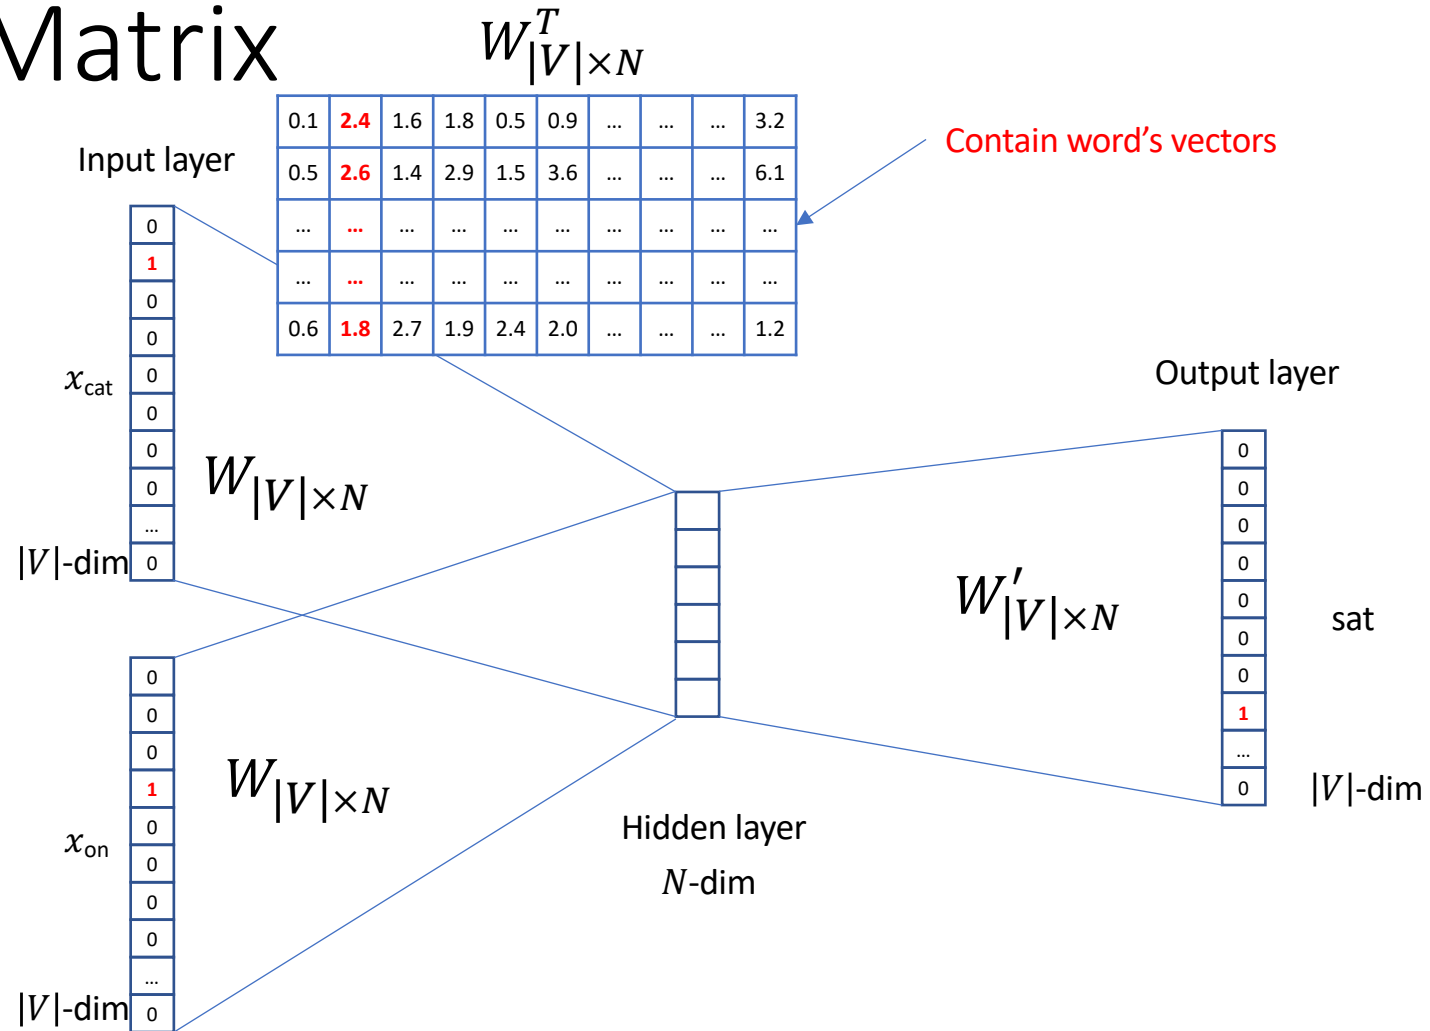


# Output





# Word's Matrix



We can consider either  $W$  or  $W'$  as the word's representation, or even take the average.

# Output: Similarity Computation

- Computed using the dot product of the two vectors
- To convert a similarity to a probability, use softmax

$$\Pr(w_k | w_{k+j}, -m \leq j \leq m, j \neq 0) = \Pr(w_k | \hat{v}) = \frac{\exp(\hat{v}^T u_k)}{\sum_{j=1}^{|V|} \exp(\hat{v}^T u_j)}$$

for all word  $w_k$

- Reminder:  $\hat{v}$  is the vector representation of the context and  $u_k$  is the vector representation of  $w_k$
- Let  $\theta$  be the set of parameters  $u_k$  for all words  $w_k$  and  $\hat{v}$  for all contexts  $\{w_{k+j}, -m \leq j \leq m, j \neq 0\}$
- In practice, use negative sampling
  - Too many words in the denominator
  - The denominator is only computed for a few words

# Word2vec: objective function

- It is the cross-entropy
- Likelihood is equivalent to cross-entropy when using one-hot encoding:

$$L(\theta) = \prod_{t=1}^T \Pr(w_t | w_{t+j}, -m \leq j \leq m, j \neq 0; \theta)$$

$$-\log L(\theta) = -\sum_{t=1}^T \log \Pr(w_t | w_{t+j}, -m \leq j \leq m, j \neq 0; \theta)$$

$$-\log L(\theta) = -\sum_{t=1}^T \sum_{x \in V} \log \Pr(w_t | w_{t+j}, -m \leq j \leq m, j \neq 0; \theta) \Pr(w_t = x)$$

# Some interesting results

## Word Analogies

Test for linear relationships, examined by Mikolov et al. (2014)

man:woman :: king:?


+ king [ 0.30 0.70 ]

- man [ 0.20 0.20 ]


+ woman [ 0.60 0.30 ]

---

queen [ 0.70 0.80 ]



# Word analogies



# Weakness of Word Embedding

- Very vulnerable, and not a robust concept
- Can take a long time to train
- Non-uniform results
- Hard to understand and visualize

# Lab: The Continuous Bag of Words (CBOW) Model

- The CBOW model architecture tries to predict the current target word (the center word) based on the source context words (surrounding words).
- Considering a simple sentence, *“the quick brown fox jumps over the lazy dog”*
  - This can be pairs of (*context\_window*, *target\_word*) where if we consider a context window of size 2, we have examples like
    - (*[quick, fox], brown*),
    - (*[the, brown], quick*),
    - (*[the, dog], lazy*)
    - and so on.
- Thus the model tries to predict the *target\_word* based on the *context\_window* words.

# Conclusion



# Conclusion

- Word embedding is a powerful preprocessing
- Used also in more advanced processing (graph processing for example)
- Embedding is not really deep learning but...
- Embedding is a very important step for deep neural network architecture dedicated to sequence analysis